



Faculty of computer science

Master of data science

Moscow, 2022

ML-based plant type classification based on time series agricultural data

Student: Elizaveta Gavrilova

Supervisors: Polina Polunina, Visiting Lecturer, Big Data and Information Retrieval

Department, Faculty of Computer Science; Specialist McKinsey and co,

Sarkis Samvelovich Grigoryan, Head of Competence center for Artificial Intelligence at

Digital Economy Development Fund,



Problem

wheat/no-wheat classification
for pre-defined NDVI



Data

- spectral data (NDVI)
- non-standard structure
- small sample due to NDA restrictions



Partner

GC «2050, DIGITAL»



Goal

find best model in terms of
classification metrics

Objectives

Dataset

Methods

Experiments

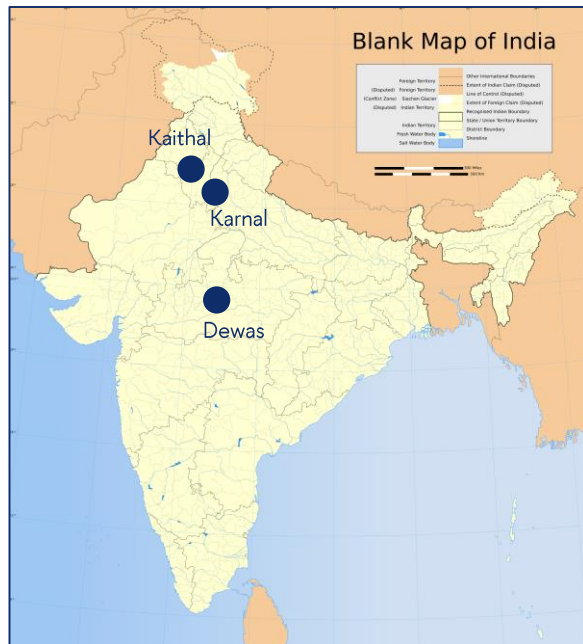
Results

Outcomes

Data description

3 Indian districts: Kaithal, Karnal and Dewas

Data points represent NDVI indices



Source: OneSoil

~200 data points
collected for each
district
Since 2020-10-20
to 2021-05-10



Source: Up42

Objectives

Dataset

Methods

Experiments

Results

Outcomes

Data description

gfid	Nº	NDVI
72001	1	0,186
72001	2	0,184
72001	3	0,183
72001
72001	200	0,185
72002	1	0,193
72002	2	0,192
72002	3	0,195
72002
72002	200	0,197

The dataset looks like the following:
At each place(gfid) NDVI index was measured 200 times. Then, all those measurements were concatenated.

So, originally the data is a 3d matrix with dimensions:

- place(gfid)
- timestamp(1 ... 200)
- NDVI

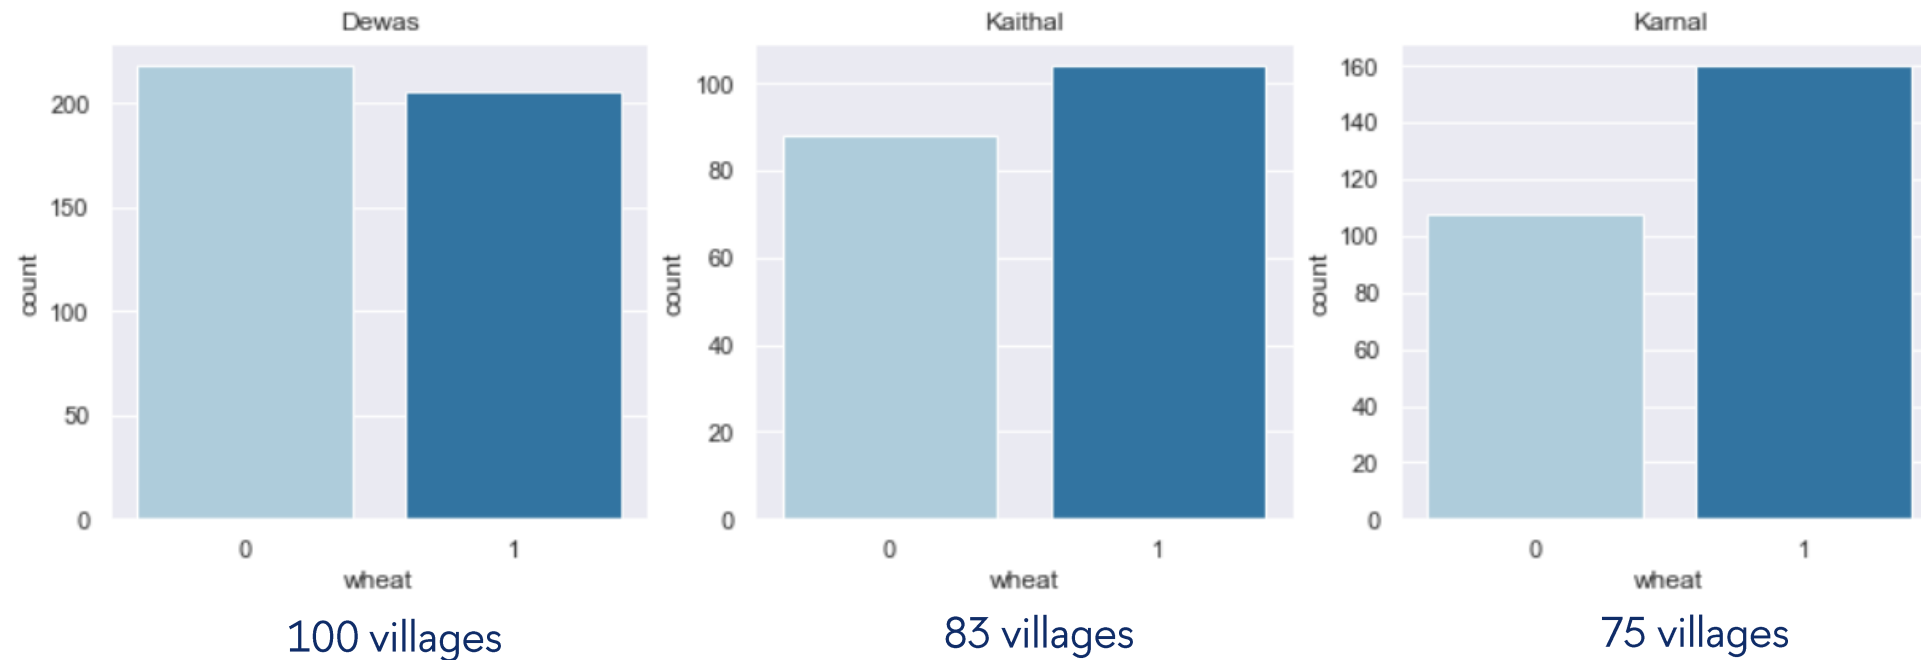
gfid	Nº	NDVI
72001	1	0,186
72001	2	0,184
72001	3	0,183
72001
72001	200	0,185
72002	1	0,193
72002	2	0,192
72002	3	0,195
72002
72002	200	0,197
...
72375	1	0,192
72375	2	0,194
72375	3	0,210
72375
72375	200	0,187



Data description

Target distribution: imbalance differs for the districts

Splitting strategy: random stratified 20/80



Objectives

Dataset

Methods

Experiments

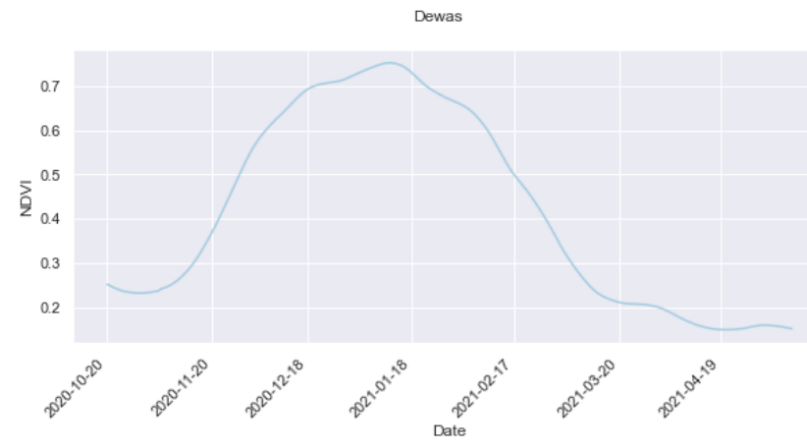
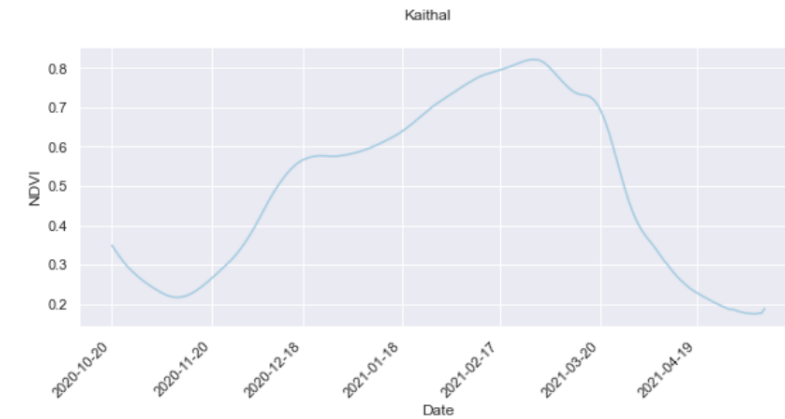
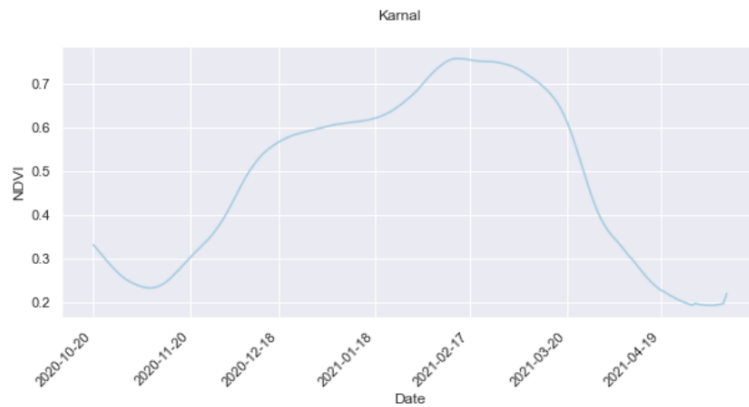
Results

Outcomes



Data description

Averaged time serieses for each district



Objectives

Dataset

Methods

Experiments

Results

Outcomes

ARIMA

Auto Regressive Integrated Moving Average (**ARIMA**) model is among one of the most popular and widely used statistical methods for time-series forecasting.

It is a class of statistical algorithms that captures the standard temporal dependencies that is unique to a time series data in order to predict future trends.

ARIMA models has 3 parameters: **q**, **d** and **p**,

- **q** represents **AR** component order,
- **d** stand for order of the integrated series,
- **p** represents **MA** component order,

Model equation is

$$\hat{y}_t = c + \theta_1 e_{t-1} + \theta_2 e_{t-2} + \dots + \theta_q e_{t-q} + e_t$$

ARIMA model assumptions:

- time series is stationary,
- residuals are homoscedastic,
- residuals are normally distributed,



ARIMA

Assumptions testing

Stationarity

Dickey-Fuller test

H_0 : there is unit root in time series
and the series is non-stationary

H_a : no unit roots in time series and
the series is stationary

Distribution: Dickey-Fuller's

Normality

D'Agostino-Pearson test

H_0 : the sample is drawn from a
normally distributed population

H_a : the sample is not drawn
from a normally distributed
population

Distribution: Chi-square

Homoscedasticity

Ljung-Box test

H_0 : the data is independently
distributed, the correlations in the
population from which the sample is
taken are 0

H_a : the data is not independently
distributed, serial correlation exists,

Distribution: Chi-square

Visual tests: QQ, Auto-correlation and Partial auto-correlation plots

Objectives

Dataset

Methods

Experiments

Results

Outcomes



Classical ML methods

Logistics Regression

Linear algorithm,
In statistics, the logistic model is a statistical model that models the probability of one event (out of two alternatives) taking place by having the log-odds for the event be a linear combination of one or more independent variables (predictors),

Random Forest

Decision trees algorithm, Random forests use a method called bagging to combine many decision trees to create an ensemble, Bagging simply means combining in parallel,

Boosting on trees

Decision trees algorithm, In boosting, new trees are formed by considering the errors of trees in previous rounds, Therefore, new trees are created one after another, Each tree is dependent on the previous tree,

SVM

SVM maps training examples to points in space so as to maximise the width of the gap between the two categories, New examples are then mapped into that same space and predicted to belong to a category based on which side of the gap they fall,

Objectives

Dataset

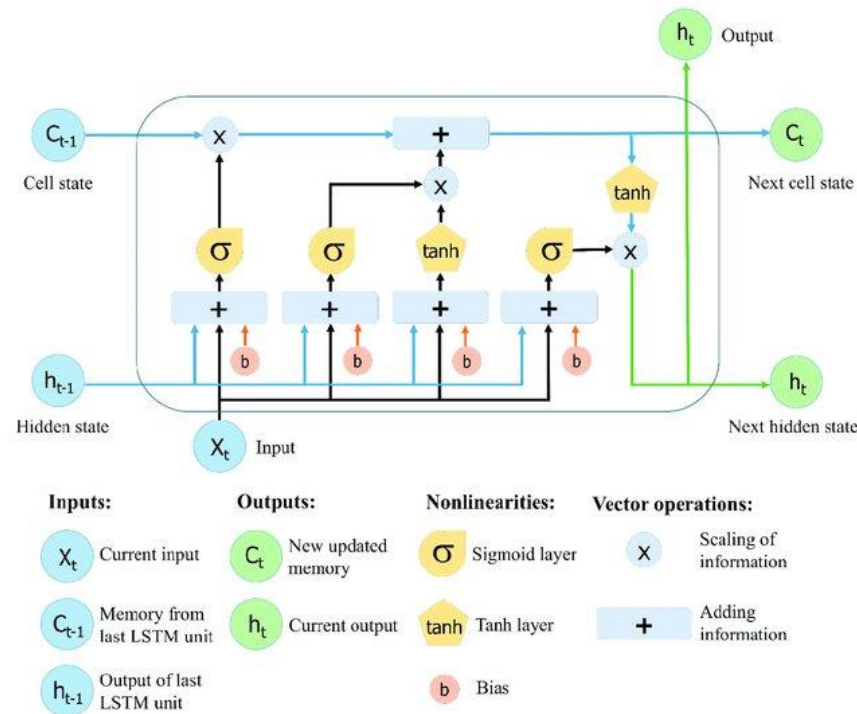
Methods

Experiments

Results

Outcomes

RNN: LSTM



Source: ResearchGate

Long Short Term Memory networks – usually just called “LSTMs” – are a special kind of RNN, capable of learning long-term dependencies,

The LSTM does have the ability to remove or add information to the cell state, carefully regulated by structures called **gates**,

Gates are a way to optionally let information through, They are composed out of a sigmoid neural net layer and a pointwise multiplication operation.

Objectives

Dataset

Methods

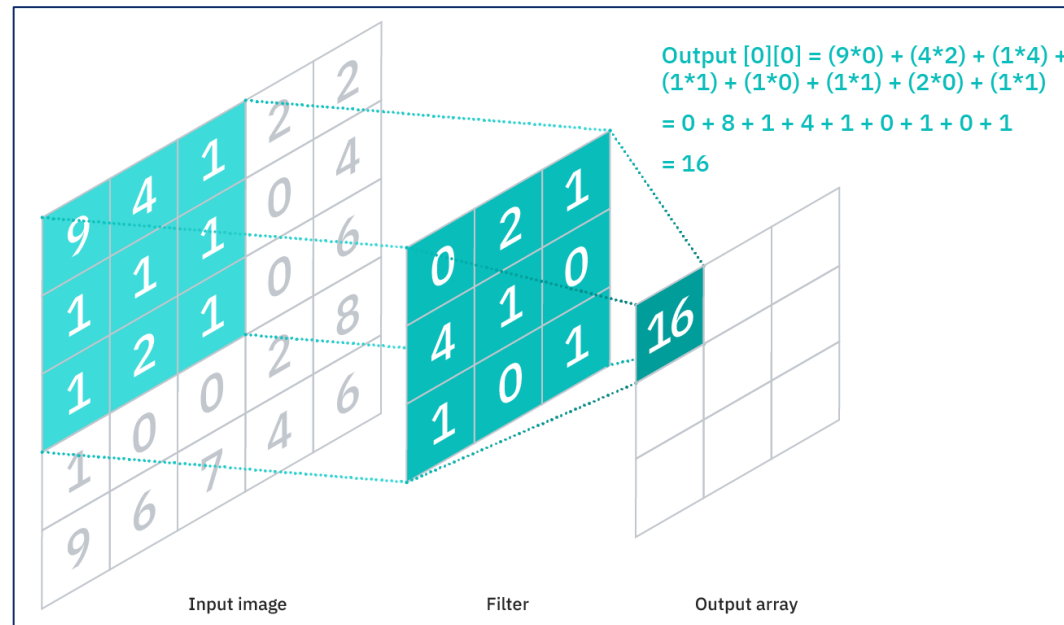
Experiments

Results

Outcomes

CNN

Convolutional neural network is composed of multiple building blocks, such as **convolution layers**, **pooling layers**, and **fully connected layers**, and is designed to automatically and adaptively learn spatial hierarchies of features.



Source: IBM

Objectives

Dataset

Methods

Experiments

Results

Outcomes

ARIMA data preparation

Given: for each data point ~200 observations

For example, at some place in Dewas with id 72001, we have a time series of collected NDVI:

	NDVI
1	0,186
2	0,184
3	0,183
...	...
200	0,185



Based on visual PAC and AC analysis and Dickey-Fuller test, parameters for Dewas district for ARIMA are $p=1$, $d=1$, $q=5$



Train ARIMA(1,1,5) for given time series and receive ARIMA model parameters to use them later as features:

	Ar,L1	Ma,L1	Ma,L2	Ma,L3	Ma,L4	Ma,L5	sigma2
72001	0,964	0,493	0,641	0,530	0,592	0,542	0,000
72002	0,837	0,351	0,578	0,555	0,563	0,521	0,004
...
72375	0,647	0,282	0,629	0,851	0,946	0,852	0,001

For Dewas, we have 375 collected time serieses

Objectives

Dataset

Methods

Experiments

Results

Outcomes

ARIMA data preparation

gfid	Nº	NDVI
72001	1	0,186
72001	2	0,184
72001	3	0,183
72001
72001	200	0,185
72002	1	0,193
72002	2	0,192
72002	3	0,195
72002
72002	200	0,197



Based on visual PAC and AC analysis and Dickey-Fuller test, parameters for Dewas district for ARIMA are $p=1$, $d=1$, $q=5$



	Ar,L1	Ma,L1	Ma,L2	Ma,L3	Ma,L4	Ma,L5	sigma2
72001	0,964	0,493	0,641	0,530	0,592	0,542	0,000
72002	0,837	0,351	0,578	0,555	0,563	0,521	0,004
...
72375	0,647	0,282	0,629	0,851	0,946	0,852	0,001

Repeat ARIMA training for each time series

For Dewas, we have 375 collected time series

Objectives

Dataset

Methods

Experiments

Results

Outcomes

Classical ML+ ARIMA

Repeat ARIMA training for each time series for each district

- Karnal: ARIMA(1,0,5)
- Kaithall: ARIMA(1,1,5)
- Dewas: ARIMA(5,1,5)



On ARIMA datasets
train:

- CatBoost
- LightGBM
- XGBoost
- RandomForest
- LogReg
- SVM



Compare results

Objectives

Dataset

Methods

Experiments

Results

Outcomes

Classical ML Results

Dewas

	F1-score	AUC-ROC
Logistic Regression	0,56	0,56
SVM	0,56	0,56
Random Forest	0,4	0,4
CatBoost	0,49	0,49
LightGBM	0,49	0,49
XGBoost	0,48	0,48

Karnal

	F1-score	AUC-ROC
Logistic Regression	0,52	0,5
SVM	0,49	0,46
Random Forest	0,56	0,52
CatBoost	0,51	0,49
LightGBM	0,54	0,49
XGBoost	0,54	0,5


Kaithal

	F1-score	AUC-ROC
Logistic Regression	0,42	0,42
SVM	0,38	0,38
Random Forest	0,5	0,49
CatBoost	0,5	0,49
LightGBM	0,44	0,43
XGBoost	0,54	0,53

CNN

Data preparation – Step 1, Add averaged NDVIs and drop NaNs created by averaging (now, start from 10ths observation)

ndvi								
1	0,186							
2	0,184							
3	0,183							
...	...							
200	0,156							



#	ndvi	ndvi_avg_2	ndvi_avg_3	...	ndvi_avg_8	ndvi_avg_9	ndvi_avg_10
10	0,184	0,184	0,185	...	0,196	0,201	0,206
11	0,185	0,185	0,184	...	0,192	0,195	0,199
12	0,187	0,186	0,185	...	0,189	0,191	0,194
13	0,189	0,188	0,187	...	0,187	0,189	0,191
14	0,191	0,190	0,189	...	0,187	0,188	0,189
...
196	0,158	0,157	0,156	...	0,152	0,151	0,151
197	0,159	0,159	0,158	...	0,153	0,153	0,152
198	0,158	0,159	0,158	...	0,155	0,154	0,153
199	0,157	0,158	0,158	...	0,156	0,155	0,154
200	0,156	0,157	0,157	...	0,156	0,156	0,155

CNN

Data preparation – Step 2, Transform dataset from 2d to 3d: crop first 150 observations into 15 smaller datasets 10x10

#	ndvi	ndvi_avg_2	ndvi_avg_3	...	ndvi_avg_8	ndvi_avg_9	ndvi_avg_10
10	0,184	0,184	0,185	...	0,196	0,201	0,206
11	0,185	0,185	0,184	...	0,192	0,195	0,199
12	0,187	0,186	0,185	...	0,189	0,191	0,194
13	0,189	0,188	0,187	...	0,187	0,189	0,191
14	0,191	0,190	0,189	...	0,187	0,188	0,189
...
196	0,158	0,157	0,156	...	0,152	0,151	0,151
197	0,159	0,159	0,158	...	0,153	0,153	0,152
198	0,158	0,159	0,158	...	0,155	0,154	0,153
199	0,157	0,158	0,158	...	0,156	0,155	0,154
200	0,156	0,157	0,157	...	0,156	0,156	0,155



	<u>ndvi</u>	ndvi_avg_2	...	ndvi_avg_10		
10	xxx	xxx	...	xxx		
11	xxx		<u>ndvi</u>	ndvi_avg_2	...	ndvi_avg_10
...	...	20	xxx	xxx	...	xxx
20	xxx	21	xxx	xxx	...	xxx
	
		30	xxx		...	xxx
		...				
			<u>ndvi</u>	ndvi_avg_2	...	ndvi_avg_10
		140	xxx	xxx	...	xxx
		141	xxx	xxx	...	xxx
	
		150	xxx	xxx	...	xxx

Objectives

Dataset

Methods

Experiments

Results

Outcomes

CNN

Data preparation – Finally: for each id we have 15 datasets,
So, the final dataset size is $n \times 15 \times 10 \times 10$, where n is a batch size.

	<u>ndvi</u>	ndvi_avg_2		...	ndvi_avg_10	
10	xxx	xxx		...	xxx	
11	xxx		<u>ndvi</u>	ndvi_avg_2	...	ndvi_avg_10
...	...	20	xxx	xxx	...	xxx
20	xxx	21	xxx	xxx	...	xxx
	
		30	xxx		...	xxx
		...				
			<u>ndvi</u>	ndvi_avg_2	...	ndvi_avg_10
		140	xxx	xxx	...	xxx
		141	xxx	xxx	...	xxx
	
		150	xxx	xxx	...	xxx



From size: $(n \times 200) \times 1$

To size: $n \times 15 \times 10 \times 10$

where n is number of ids
for training (or batch size)

Objectives

Dataset


Methods

Experiments

Results

Outcomes

CNN architecture



Objectives

Dataset

Methods

Experiments

Results

Outcomes



Deep Learning Results

Dewas

	F1-score	AUC-ROC
LSTM	0,42	0,47
CNN	0,8	0,8

Karnal

	F1-score	AUC-ROC
LSTM	0,61	0,56
CNN	0,76	0,74

Kaithal

	F1-score	AUC-ROC
LSTM	0,42	0,57
CNN	0,56	0,55

Objectives

Methods

Dataset

Experiments

Results

Outcomes

Outcomes

CNN is the best solution because:

- F1-score is 15-30% higher for the CNN
- Recall does not vary heavily for different classes
- Lack of time-extensive feature engineering step for CNN like ARIMA for classic ML
- Lower deviation in results for different districts

-
- The partner «2050, DIGITAL» is interested in non-standard feature engineering technique like ARIMA modelling
 - The CNN metric results exceeded the partner's expectations



Further steps

Business:

- Add business logic to the model to create an application
- Wrap the code with any container like Docker, orcheater with K8s
- Add web interface for users
- Add train automatizing flow like AirFlow or KubeFlow

Research:

- Tune hypereparameters for the CNN individually for each district
- Add ARIMA features to CNN



Thank you for you attention!

Objectives

Methods

Dataset

Experiments

Results

Outcomes